

Announcements

- HW1 has been due yesterday at 6 pm
- HW2 will be out later today
 - 2 weeks for completing it
 - About randomized algorithms and DP
- Fellow student Yasunari built a Discord Server to facilitate class discussion, study groups, knowing your classmates, etc.
 - Link: <https://discord.gg/JPp8gMZ>
 - Also see his piazza post:
<https://piazza.com/class/ke5yn1pnhg61l?cid=16>

CS6161: Design and Analysis of Algorithms (Fall 2020)

Greedy Algorithms

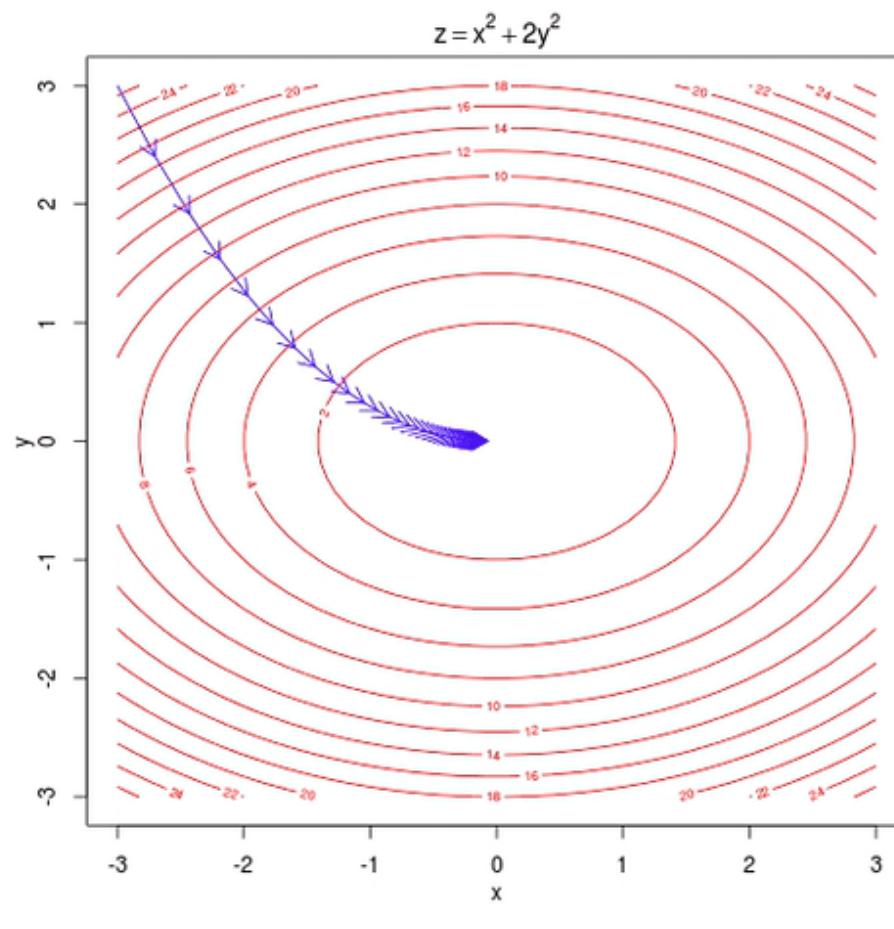
Instructor: Haifeng Xu

Outline

- Greedy Algorithms and a Simple Example
- Greedy Algorithm Analysis for a Non-trivial Problem

What is a Greedy Algorithm?

- Build up a solution by greedily **improve local or current state**, and **hope** that it will lead to a globally optimal solution



What is a Greedy Algorithm?

Pros:

- Easy to implement and often run fast
- Can be tried pretty much for any problem
 - Often leads to first-cut heuristic when problems are not well understood
- If works, typically imply useful structures of the problem

“Greed ... is good. Greed is right. Greed works.”

-- Michael Douglas, *Wall Street* (1987)

Unfortunately, not true in algorithm design ...

What is a Greedy Algorithm?

Pros:

- Easy to implement and often run fast
- Can be tried pretty much for any problem
 - Often leads to first-cut heuristic when problems are not well understood
- If works, typically imply useful structures of the problem

Cons:

- **Very often**, greedy algorithms do not work
- Easy to invent many versions of greedy algorithms for a problem, but finding the right one and proving its correctness is not easy

Unlike DP and D-and-C, greedy algorithms require a proof of correctness in most cases

A Warm-up Example

- A worker has total work time T
- There are n jobs to be completed
 - Job i takes time t_i to complete and results in payoff r_i
 - May complete any α fraction of job i , taking αt_i time with payoff αr_i

Algorithmic Question

Compute the fraction α_i of job i for all i , to maximize payoff

Is this a knapsack problem?

- Not exactly, in knapsack cannot pack 0.5 fraction of an item
- It's called **fractional knapsack**, and is much easier

Fractional Knapsack – Greed Works

Now you can complete any fraction of a job and getting payoff proportional to it, which job is your most profitable choice?

- The one with maximum payoff r_i ?
- The one with shortest time t_i ?
- Should be the one with largest average payoff $\frac{r_i}{t_i}$!

Fact: within time t where $t \leq t_i$ for i , doing job i^* with maximum $\frac{r_i}{t_i}$ achieves the largest payoff.

Proof: Payoff of completing job i is

$$\frac{t}{t_i} \times r_i$$

 fraction of job i completed

Fractional Knapsack – Greed Works

Now you can complete any fraction of a job and getting payoff proportional to it, which job is your most profitable choice?

- The one with maximum payoff r_i ?
- The one with shortest time t_i ?
- Should be the one with largest average payoff $\frac{r_i}{t_i}$!

Fact: within time t where $t \leq t_i$ for i , doing job i^* with maximum $\frac{r_i}{t_i}$ achieves the largest payoff.

Proof: Payoff of completing job i is

$$\frac{t}{t_i} \times r_i = t \times \frac{r_i}{t_i}$$



fraction of job i completed

Fractional Knapsack – Greed Works

Now you can complete any fraction of a job and getting payoff proportional to it, which job is your most profitable choice?

- The one with maximum payoff r_i ?
- The one with shortest time t_i ?
- Should be the one with largest average payoff $\frac{r_i}{t_i}$!

Fact: within time t where $t \leq t_i$ for i , doing job i^* with maximum $\frac{r_i}{t_i}$ achieves the largest payoff.

Proof: Payoff of completing job i is

$$\frac{t}{t_i} \times r_i = t \times \frac{r_i}{t_i} \quad \text{which is maximized at } i^*$$

↑
fraction of job i completed

Fractional Knapsack – Greed Works

Now you can complete any fraction of a job and getting payoff proportional to it, which job is your most profitable choice?

- The one with maximum payoff r_i ?
- The one with shortest time t_i ?
- Should be the one with largest average payoff $\frac{r_i}{t_i}$!

Fact: within time t where $t \leq t_i$ for i , doing job i^* with maximum $\frac{r_i}{t_i}$ achieves the largest payoff.

Remark: this fact is the reason that Greed works.

Fractional Knapsack: Pseudo-Code

FractionalKnapsack-Greed($T, \{t_i, r_i\}_{i=1}^n$)

1

2

3

4

5

6

7

8

9

10

Fractional Knapsack: Pseudo-Code

FractionalKnapsack-Greed($T, \{t_i, r_i\}_{i=1}^n$)

1 Sort jobs based on $\frac{r_i}{t_i}$ with descending order

2

3

4

5

6

7

8

9

10

Fractional Knapsack: Pseudo-Code

FractionalKnapsack-Greed($T, \{t_i, r_i\}_{i=1}^n$)

- 1 Sort jobs based on $\frac{r_i}{t_i}$ with descending order
- 2 $t = T; R = 0$ // t : total remaining time; R : total rewards so far
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10

Fractional Knapsack: Pseudo-Code

FractionalKnapsack-Greed($T, \{t_i, r_i\}_{i=1}^n$)

- 1 Sort jobs based on $\frac{r_i}{t_i}$ with descending order
- 2 $t = T; R = 0$ // t : total remaining time; R : total rewards so far
- 3 **for** $i = 1, \dots, n$ in the sorted order
- 4 **if** $t > t_i$
- 5 $R = R + r_i$
- 6 $t = t - t_i$
- 7
- 8
- 9
- 10

Fractional Knapsack: Pseudo-Code

FractionalKnapsack-Greed($T, \{t_i, r_i\}_{i=1}^n$)

```
1  Sort jobs based on  $\frac{r_i}{t_i}$  with descending order
2   $t = T; R = 0$  //  $t$ : total remaining time;  $R$ : total rewards so far
3  for  $i = 1, \dots, n$  in the sorted order
4      if  $t > t_i$ 
5           $R = R + r_i$ 
6           $t = t - t_i$ 
7      else
8           $R = R + r_i \times \frac{t}{t_i}$ 
9      return  $R$ 
10
```

Fractional Knapsack: Pseudo-Code

FractionalKnapsack-Greed($T, \{t_i, r_i\}_{i=1}^n$)

```
1  Sort jobs based on  $\frac{r_i}{t_i}$  with descending order
2   $t = T; R = 0$  //  $t$ : total remaining time;  $R$ : total rewards so far
3  for  $i = 1, \dots, n$  in the sorted order
4      if  $t > t_i$ 
5           $R = R + r_i$ 
6           $t = t - t_i$ 
7      else
8           $R = R + r_i \times \frac{t}{t_i}$ 
9      return  $R$ 
10     return  $R$ 
```

Fractional Knapsack: Pseudo-Code

FractionalKnapsack-Greed($T, \{t_i, r_i\}_{i=1}^n$)

- 1 Sort jobs based on $\frac{r_i}{t_i}$ with descending order
- 2 $t = T; R = 0$ // t : total remaining time; R : total rewards so far
- 3 **for** $i = 1, \dots, n$ in the sorted order
- 4 **if** $t > t_i$
- 5 $R = R + r_i$
- 6 $t = t - t_i$
- 7 **else**
- 8 $R = R + r_i \times \frac{t}{t_i}$
- 9 **return** R
- 10 **return** R

$O(n \log n)$ time due to sorting

Fractional Knapsack: Pseudo-Code

FractionalKnapsack-Greed($T, \{t_i, r_i\}_{i=1}^n$)

```
1  Sort jobs based on  $\frac{r_i}{t_i}$  with descending order
2   $t = T; R = 0$  //  $t$ : total remaining time;  $R$ : total rewards so far
3  for  $i = 1, \dots, n$  in the sorted order
4      if  $t > t_i$ 
5           $R = R + r_i$ 
6           $t = t - t_i$ 
7      else
8           $R = R + r_i \times \frac{t}{t_i}$ 
9      return  $R$ 
10 return  $R$ 
```

$O(n \log n)$ time due to sorting

Remark: there is a better $O(n)$ implementation using the *median of median algorithm* (we did not cover though) that can avoid sorting. Think about it if you are interested!

Correctness is Relatively Easy Here

Claim: at any time t , the collected reward so far is the largest possible.

Follows from previous fact, as we are completing the jobs in descending order of $\frac{r_i}{t_i}$.

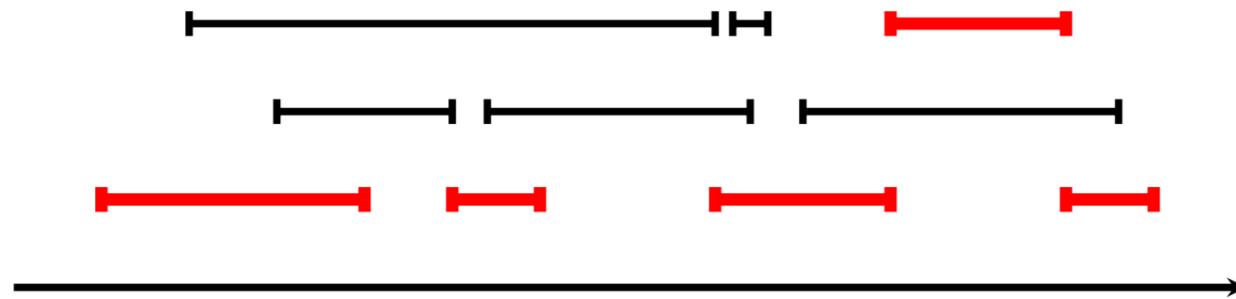
```
FractionalKnapsack-Greed( $T, \{t_i, r_i\}_{i=1}^n$ )
1   Sort jobs based on  $\frac{r_i}{t_i}$  with descending order
2    $t = T; R = 0$ 
3   for  $i = 1, \dots, n$  in the sorted order
4       if  $t > t_i$ 
5            $R = R + r_i$ 
6            $t = t - t_i$ 
7       else
8            $R = R + r_i \times \frac{t}{t_i}$ 
9       return  $R$ 
10  return  $R$ 
```

Outline

- Greedy Algorithms and a Simple Example
- Greedy Algorithm Analysis for a Non-trivial Problem

Example 2: Interval Job Scheduling

- **Input:** a set of jobs, job i starts at s_i and finishes at f_i
 - Two jobs are compatible if they do not overlap
- **Output:** find a maximum subset of mutually compatible jobs



The Greedy Framework

- Okay, so we know it's going to be some greedy algorithm...

```
SomeGreedyAlgo( $\{s_i, f_i\}_{i=1}^n$ )
```

```
1
```

```
2
```

```
3
```

```
4
```

```
5
```

```
6
```

```
7
```

The Greedy Framework

➤ Okay, so we know it's going to be some greedy algorithm...

```
SomeGreedyAlgo( $\{s_i, f_i\}_{i=1}^n$ )
```

- 1 Let R be the set of all jobs;

- 2 Initialize $X = 0$ // X will store all scheduled jobs

- 3

- 4

- 5

- 6

- 7

The Greedy Framework

➤ Okay, so we know it's going to be some greedy algorithm...

```
SomeGreedyAlgo( $\{s_i, f_i\}_{i=1}^n$ )
```

- 1 Let R be the set of all jobs;
- 2 Initialize $X = 0$ // X will store all scheduled jobs
- 3 **while** R is not empty
- 4 choose some $i \in R$;
- 5 add i to X
- 6
- 7

The Greedy Framework

➤ Okay, so we know it's going to be some greedy algorithm...

```
SomeGreedyAlgo( $\{s_i, f_i\}_{i=1}^n$ )
```

- 1 Let R be the set of all jobs;
- 2 Initialize $X = 0$ // X will store all scheduled jobs
- 3 **while** R is not empty
- 4 choose some $i \in R$;
- 5 add i to X
- 6 remove from R all jobs that overlap with i
- 7 **return** X

The Greedy Framework

➤ Okay, so we know it's going to be some greedy algorithm...

```
SomeGreedyAlgo( $\{s_i, f_i\}_{i=1}^n$ )
```

- 1 Let R be the set of all jobs;
- 2 Initialize $X = 0$ // X will store all scheduled jobs
- 3 **while** R is not empty
- 4 choose some $i \in R$;
- 5 add i to X
- 6 remove from R all jobs that overlap with i
- 7 **return** X

Tricky part: decide which job from R to process

The Greedy Framework

Candidate choice of next $i \in R$

- **Earliest start time:** the $i \in R$ with smallest s_i
- **Earliest finish time:** the $i \in R$ with smallest f_i
- **Shortest interval time:** the $i \in R$ with smallest $f_i - s_i$
- **Fewest conflicts:** the $i \in R$ with minimum #conflicts with other jobs

Does any of these work, and if any, which one?

Does Earliest Start Time Work?

➤ No

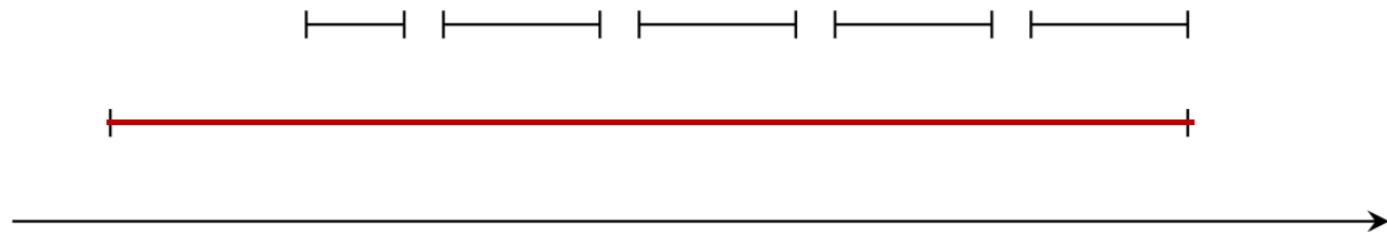


Figure: counterexample for earliest start time

Does Shortest Interval Choice Work?

➤ No

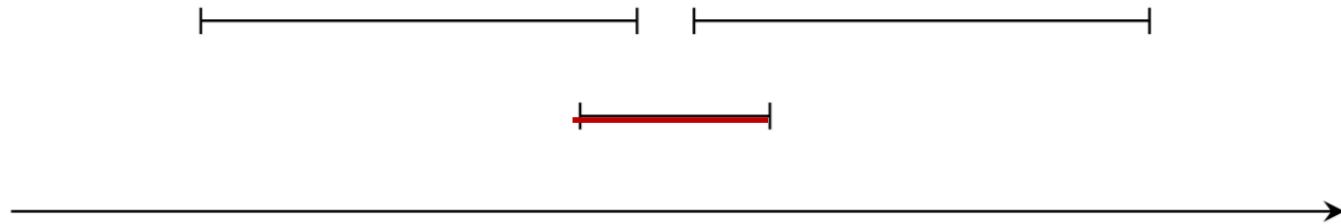


Figure: counterexample for shortest interval

Does Fewest Conflicts Choice Work?

➤ No

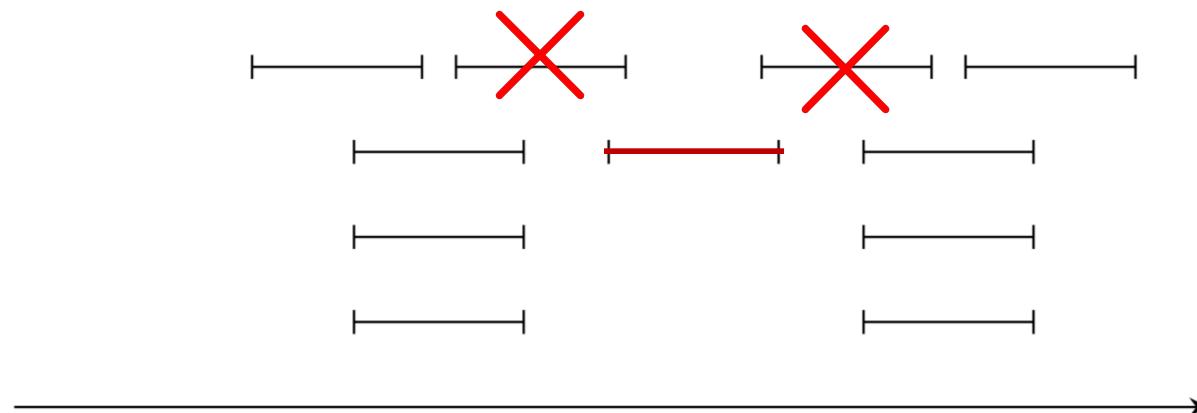


Figure: counterexample for fewest conflicts

Does Fewest Conflicts Choice Work?

➤ No

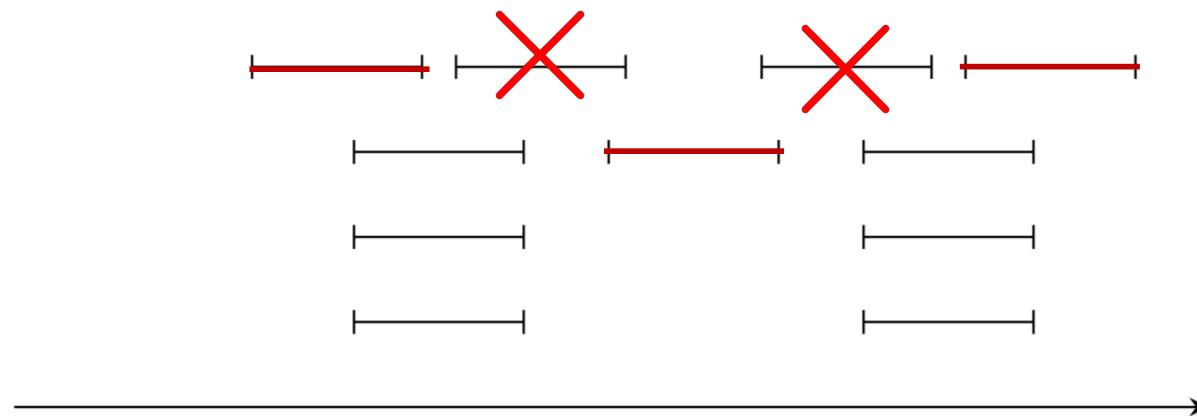


Figure: counterexample for fewest conflicts

Remaining Choice: Earliest Finish Time

- Fortunately, this turns out to work

```
GreedyAlgo( $\{s_i, f_i\}_{i=1}^n$ )
```

- 1 Let R be the set of all jobs;
- 2 Initialize $X = 0$ // X will store all scheduled jobs
- 3 **while** R is not empty
- 4 choose $i \in R$ with smallest f_i ;
- 5 add i to X
- 6 remove from R all jobs that overlap with i
- 7 **return** X

- Need to sort jobs in R according to f_i every time
- Better idea: sort all jobs at the very beginning

Remaining Choice: Earliest Finish Time

- Fortunately, this turns out to work

```
GreedyAlgo( $\{s_i, f_i\}_{i=1}^n$ )
```

```
1 Let  $R$  = set of all jobs sorted increasingly on  $f_i$ ;  
2 Initialize  $X = 0$  //  $X$  will store all scheduled jobs  
3 while  $R$  is not empty  
4     choose  $i \in R$  with smallest  $f_i$ ;  
5     add  $i$  to  $X$   
6     remove from  $R$  all jobs that overlap with  $i$   
7 return  $X$ 
```

- Need to sort jobs in R according to f_i every time
- Better idea: sort all jobs at the very beginning

Running Time Analysis

GreedyAlgo($\{s_i, f_i\}_{i=1}^n$)

- 1 Let R = set of all jobs sorted increasingly on f_i ;
- 2 Initialize $X = 0$ // X will store all scheduled jobs
- 3 **while** R is not empty
- 4 choose $i \in R$ with smallest f_i ; $\longrightarrow O(1)$
- 5 add i to X
- 6 remove from R all jobs that overlap with i
- 7 **return** X

$O(n \log n)$

Naïve implementation $O(n)$:
check jobs one by one

In total: $O(n^2)$

Running Time Analysis

GreedyAlgo($\{s_i, f_i\}_{i=1}^n$)

- 1 Let R = set of all jobs sorted increasingly on f_i ;
- 2 Initialize $X = 0$ // X will store all scheduled jobs
- 3 **while** R is not empty
- 4 choose $i \in R$ with smallest f_i ;
- 5 add i to X
- 6 remove from R all jobs that overlap with i
- 7 **return** X

- Step 6 has a faster implementation where each job is visited once
 - Which job j would overlap with our choice i ?
 - Exactly those with $s_j < f_i$ (since our choice implies $f_j \geq f_i$)
 - Can have another job order sorted according to s_j at the beginning; Step 6 simply removes all jobs before our chosen i in that order
- Running time under this implementation: $O(n \log n)$

Correctness Analysis

Theorem: The Greedy algorithm that picks the job with earliest finish time at each step is optimal.

- First: **feasibility** – does the algorithm output (at least) a valid solution?
 - Yes, by definition, as we removed all jobs conflicting with current schedules

Correctness Analysis

Theorem: The Greedy algorithm that picks the job with earliest finish time at each step is optimal.

- Second: **optimality** – is the output the best possible?

Correctness Analysis

Theorem: The Greedy algorithm that picks the job with earliest finish time at each step is optimal.

➤ How to prove optimality?

- Let the set of jobs achieves optimality be $O = \{o_1, \dots, o_m\}$
- Let our output set $X = \{x_1, \dots, x_k\}$

Prove $O = X$? \rightarrow Not likely

Correctness Analysis

Theorem: The Greedy algorithm that picks the job with earliest finish time at each step is optimal.

➤ How to prove optimality?

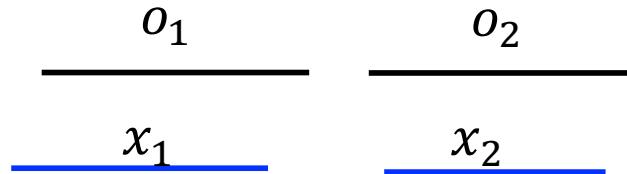
- Let the set of jobs achieves optimality be $O = \{o_1, \dots, o_m\}$
- Let our output set $X = \{x_1, \dots, x_k\}$ (**note $k \leq m$**)
- Want to prove $k = m$, i.e., our output is equally good

Correctness Analysis

Theorem: The Greedy algorithm that picks the job with earliest finish time at each step is optimal.

➤ How to prove optimality?

- Let the set of jobs achieves optimality be $O = \{o_1, \dots, o_m\}$
- Let our output set $X = \{x_1, \dots, x_k\}$ (**note $k \leq m$**)
- Want to prove $k = m$, i.e., our output is equally good
- **Claim:** for $i = 1, \dots, k$, replacing o_i in optimal O by x_i maintain optimality and feasibility
 - ❖ Induction: when $i = 1$, this is true as x_1 finishes before o_1 and thus not conflicts with o_2

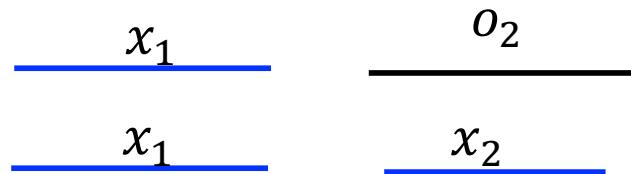


Correctness Analysis

Theorem: The Greedy algorithm that picks the job with earliest finish time at each step is optimal.

➤ How to prove optimality?

- Let the set of jobs achieves optimality be $O = \{o_1, \dots, o_m\}$
- Let our output set $X = \{x_1, \dots, x_k\}$ (**note $k \leq m$**)
- Want to prove $k = m$, i.e., our output is equally good
- **Claim:** for $i = 1, \dots, k$, replacing o_i in optimal O by x_i maintain optimality and feasibility
 - ❖ Induction: when $i = 1$, this is true as x_1 finishes before o_1 and thus not conflicts with o_2

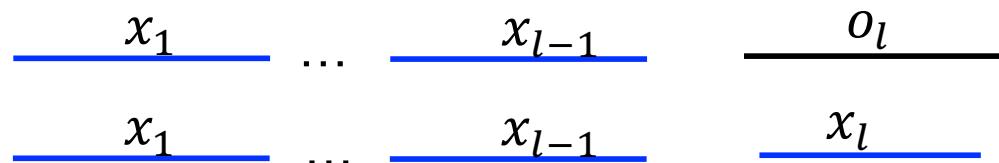


Correctness Analysis

Theorem: The Greedy algorithm that picks the job with earliest finish time at each step is optimal.

➤ How to prove optimality?

- Let the set of jobs achieves optimality be $O = \{o_1, \dots, o_m\}$
- Let our output set $X = \{x_1, \dots, x_k\}$ (**note $k \leq m$**)
- Want to prove $k = m$, i.e., our output is equally good
- **Claim:** for $i = 1, \dots, k$, replacing o_i in optimal O by x_i maintain optimality and feasibility
 - ❖ Induction: when $i = 1$, this is true as x_1 finishes before o_1 and thus not conflicts with o_2
 - ❖ After replacement for all $i \leq l - 1$, can also do the replacement for $i = l$ because x_l has earliest finish time among all remaining non-conflicting jobs



Correctness Analysis

Theorem: The Greedy algorithm that picks the job with earliest finish time at each step is optimal.

➤ How to prove optimality?

- Let the set of jobs achieves optimality be $O = \{o_1, \dots, o_m\}$
- Let our output set $X = \{x_1, \dots, x_k\}$ (**note $k \leq m$**)
- Want to prove $k = m$, i.e., our output is equally good
- **Claim: for $i = 1, \dots, k$, replacing o_i in optimal O by x_i maintain optimality and feasibility**
 - ❖ Induction: when $i = 1$, this is true as x_1 finishes before o_1 and thus not conflicts with o_2
 - ❖ After replacement for all $i \leq l - 1$, can also do the replacement for $i = l$ because x_l has earliest finish time among all remaining non-conflicting jobs
 - ❖ Since we can do this for any $i = 1, \dots, k$, at the end the optimal solution becomes the same as X
 - ❖ So $k = m$

More About Interval Scheduling

- Question 1: the greedy algorithm of picking the shortest interval instead picks at least half the optimal number of jobs.
- Question 2: what if each job i has a reward r_i , you want to schedule jobs to maximize total rewards, without conflicts
 - The case we studied is $r_i = 1$ for all i
 - This will be in your HW 2
- Question 3: what if you do not know the jobs in advance; instead, they arrive in an **online** fashion
 - An active research area

Next Lectures: Graph Algorithms

Thank You

Haifeng Xu

University of Virginia

hx4ad@virginia.edu